

Build reliable, traceable, distributed systems with ZeroMQ

PYCON2012

@jpetazzo
@dot_cloud



dotCloud

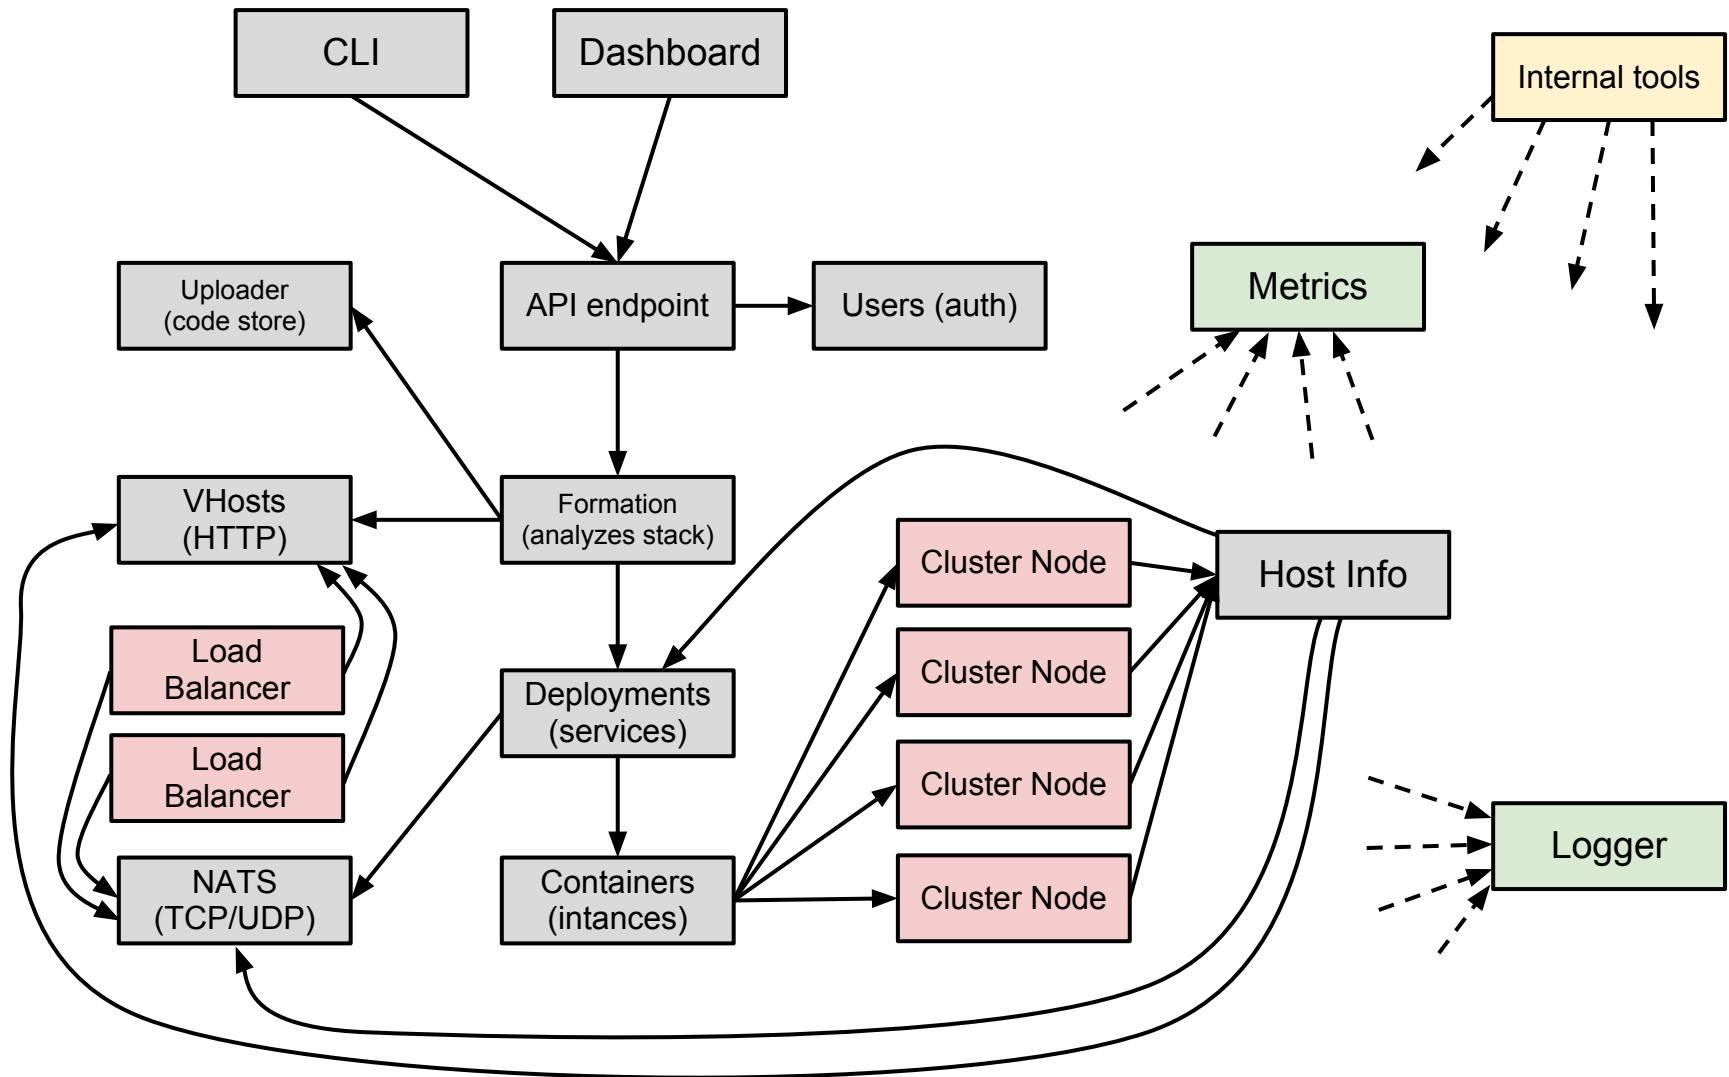
Outline

- Who we are
- Quick tour of ZeroRPC features
- Implementation overview

Introduction: why?

- We are running a PAAS:
we deploy, monitor, and scale your apps
(in the Cloud!)
- Many moving parts
- ... On a large distributed cluster

dotCloud architecture (simplified) — components, **servers**, **things receiving data from everything else**, **me**



Our requirements (easy)

- We want to be able to expose arbitrary code, with minimal modification (if any)

If we can do **import foo; foo.bar(42)**,

we want to be able to do **foo=RemoteService(...); foo.bar(42)**

- We want a self-documented system

We want to see methods, signatures, docstrings, without opening the code of the remote service, or relying on (always out-dated) documentation

Our requirements (harder)

- We want to propagate exceptions properly
- We want to be language-agnostic
- We want to be brokerless, highly available, fast, and support fan-in/fan-out topologies
(Not necessarily everything at the same time!)
- We want to trace & profile nested calls

Which call initiated the subcall which raised **SomeException**?
Which subcall of this complex call causes it to take forever?

Why not {x} ?

- **x=HTTP**

- too much overhead for some use-cases (logs and metrics)
- continuous stream of updates requires extra work (chunked encoding, websockets...)
- asynchronous notification requires extra work (long polling, websockets...)
- fan-out topology requires extra work (implementing some kind of message queue)

- **x=AMQP (or similar)**

- too much overhead for some use-cases
- requires a broker from the ground up

What we came up with

ZeroRPC!

- Based on ZeroMQ and MessagePack
- Supports everything we needed!
- Multiple implementations
 - Internal "reference" implementation in Python
 - Public "alternative" implementation with gevents
 - Internal Node.js implementation (so-so)

Example: unmodified code

Expose the "urllib" module over RPC:

```
$ zerorpc-client --server --bind tcp://0:1234 urllib
```

Yup. That's it.

There is now something listening on TCP port 1234, and exposing the Python "urllib" module.

Example: calling code

From the command line (for testing):

```
$ zerorpc-client tcp://0:1234 quote "hello pycon"  
"hello%20pycon"
```

From Python code:

```
>>> import zerorpc  
>>> remote_urllib = zerorpc.Client( )  
>>> remote_urllib.connect('tcp://0:1234')  
>>> remote_urllib.quote('hello pycon')  
'hello%20pycon'
```

Example: introspection

We can list methods:

```
$ zerorpc-client tcp://localhost:1234 | grep ^q
```

```
quote                quote('abc def') -> 'abc%20def'
```

```
quote_plus          Quote the query fragment of a URL; replacing ' ' with '+'
```

We can see signatures and docstrings:

```
$ zerorpc-client tcp://localhost:1234 quote_plus -?
```

```
quote_plus(s, safe="")
```

```
Quote the query fragment of a URL; replacing ' ' with '+'
```

Example: exceptions

\$ zerorpc-client tcp://localhost:1234 quote_plus

Traceback (most recent call last):

**File "/home/jpetazzo/.virtualenvs/dotcloud_develop/bin/zerorpc-client", line 131, in <module>
main()**

**File "/home/jpetazzo/.virtualenvs/dotcloud_develop/bin/zerorpc-client", line 127, in main
pprint(client(args.command, *parameters))**

**File "/home/jpetazzo/Work/DOTCLOUD/dotcloud/zerorpc/zerorpc.py", line 362, in __call__
return self.recv()**

**File "/home/jpetazzo/Work/DOTCLOUD/dotcloud/zerorpc/zerorpc.py", line 243, in recv
ret = self.handler(headers, method, *args)**

**File "/home/jpetazzo/Work/DOTCLOUD/dotcloud/zerorpc/zerorpc.py", line 323, in handle_result
raise e**

TypeError: quote() takes at least 1 argument (0 given)

Example: load balancing

Start a load balancing hub:

```
$ cat foo.yml  
in: "tcp://*:1111"  
out: "tcp://*:2222"  
type: queue  
$ zerohub.py foo.yml
```

Start (at least) one worker:

```
$ zerorpc-client --server tcp://localhost:2222 urllib
```

Now connect to the "in" side of the hub:

```
$ zerorpc-client tcp://localhost:1111
```

Example: high availability

Start a local HAProxy in TCP mode, dispatching requests to 2 or more remote services or hubs:

```
$ cat haproxy.cfg  
listen zerorpc 0.0.0.0:1111  
    mode tcp  
    server backend_a localhost:2222 check  
    server backend_b localhost:3333 check  
$ haproxy -f haproxy.cfg
```

Start (at least) one backend:

```
$ zerorpc-client --server --bind tcp://0:2222 urllib
```

Now connect to HAProxy:

```
$ zerorpc-client tcp://localhost:1111
```

Non-Example: PUB/SUB (not in public repo—*yet*)

- Broadcast a message to a group of nodes
 - But if a node leaves and rejoins, he'll lose messages
- Send a continuous stream of information
 - But if a speaker or listener leaves and rejoins...

You generally don't want to do this!

Better pattern: ZeroRPC streaming with
gevents

Example: streaming

- Server code returns an ~~list~~ iterator
- Client code gets an ~~list~~ iterator
- Small messages, high latency? No problem!
 - Server code will pre-push elements
 - Client code will notify server if pipeline runs low
- Huge messages? No problem!
 - Big data sets can be nicely chunked
 - They don't have to fit entirely in memory
 - Don't worry about timeouts anymore
- Also supports long polling

Example: tracing (not in public repo—yet)

```
$ dotcloud --trace alias add sushiblog.web www.deliciousrawfish.com
```

```
TraceID: 48aca4f4-75d5-40f2-b5bd-73c40ca40980
```

Ok. Now please add the following DNS record:

www.deliciousrawfish.com. IN CNAME gateway.dotcloud.com.

```
$ dotcloud-tracedump 48aca4f4-75d5-40f2-b5bd-73c40ca40980
```

```
[2012-03-07 23:56:17.759738] uwsgi@api --- run_command ---> api_00@zeroworkers (lag: 1ms | exec: 98ms)
[2012-03-07 23:56:17.770432] api_00@zeroworkers --- track ---> mixpanel_00@zeroworkers (lag: 1ms | exec: 70ms)
[2012-03-07 23:56:17.771264] api_00@zeroworkers --- km_track ---> mixpanel_01@zeroworkers (lag: 1ms | exec: 71ms)
[2012-03-07 23:56:17.771994] api_00@zeroworkers --- km_track ---> mixpanel_02@zeroworkers (lag: 1ms | exec: 71ms)
[2012-03-07 23:56:17.773041] api_00@zeroworkers --- record_event ---> users-events@users (lag: 0ms | exec: 0ms)
[2012-03-07 23:56:17.774972] api_00@zeroworkers --- info (pending) ---> formation-in [Not Received]
[2012-03-07 23:56:17.783590] api_00@zeroworkers --- info (pending) ---> formation-in [Not Received]
[2012-03-07 23:56:17.830107] api_00@zeroworkers --- add_alias ---> deployments_04@deployments (lag: 0ms | exec: 27ms)
[2012-03-07 23:56:17.831453] deployments_04@deployments --- exists_frontend ---> vhosts@vhosts (lag: 0ms | exec: 4ms)
[2012-03-07 23:56:17.836288] vhosts@vhosts --- OK ---> deployments_04@deployments (lag: 0ms | exec: 0ms)
[2012-03-07 23:56:17.837370] deployments_04@deployments --- exists ---> vhosts@vhosts (lag: 0ms | exec: 1ms)
[2012-03-07 23:56:17.838733] vhosts@vhosts --- OK ---> deployments_04@deployments (lag: 0ms | exec: 0ms)
[2012-03-07 23:56:17.840068] deployments_04@deployments --- add_frontend ---> vhosts@vhosts (lag: 0ms | exec: 15ms)
[2012-03-07 23:56:17.856166] vhosts@vhosts --- OK ---> deployments_04@deployments (lag: 0ms | exec: 0ms)
[2012-03-07 23:56:17.857647] deployments_04@deployments --- OK ---> api_00@zeroworkers (lag: 0ms | exec: 0ms)
[2012-03-07 23:56:17.859824] api_00@zeroworkers --- OK ---> uwsgi@api (lag: 0ms | exec: 0ms)
```

Implementation details

This will be useful if...

- You think you might want to use ZeroRPC
- You think you might want to hack ZeroRPC
- You want to reimplement something similar
- You just happen to love distributed systems

ØMQ

- Sockets on steroids
<http://zguide.zeromq.org/page:all>
- Handles (re)connections for us
- Works over regular TCP
- Also has superfast ipc:// and inproc://
- Different patterns:
 - REQ/REP (basic, synchronous RPC call + response)
 - PUB/SUB (shout and listen to streams of events)
 - PUSH/PULL (load balance or collect messages)
 - DEALER/ROUTER (REQ/REP with routing)
- **pip install pyzmq-static FTW** (thanks @brandon_rhodes!)

Serialization: MessagePack

In our tests, msgpack is more efficient than JSON, BSON, YAML:

- 20-50x faster
- serialized output is 2x smaller or better

\$ pip install msgpack-python

```
>>> import msgpack
```

```
>>> bytes = msgpack.dumps(data)
```

Wire format

Request: (headers, method_name, args)

- headers dict
 - no mandatory header
 - carries the protocol version number
 - used for tracing in our in-house version
- args
 - list of arguments
 - no named parameters

Response: (headers, ERR|OK|STREAM, value)

Timeouts

- ØMQ does not detect disconnections (or rather, it works hard to hide them)
- You can't know when the remote is gone
- Original implementation: 30s timeout
- Published implementation: heartbeat

Introspection

- Expose a few special calls:
 - `_zerorpc_list` to list calls
 - `_zerorpc_name` to know who you're talking to
 - `_zerorpc_ping` (redundant with the previous one)
 - `_zerorpc_help` to retrieve the docstring of a call
 - `_zerorpc_args` to retrieve the argspec of a call
 - `_zerorpc_inspect` to retrieve everything at once
- Introspection + service discovery = WIN

Naming

- Published implementation does not include any kind of naming/discovery
- In-house version uses a flat YAML file, mapping service names to ØMQ addresses and socket types
- In progress: use DNS records
 - SRV for host+port
 - TXT for ØMQ socket type (not sure about this!)
- In progress: registration of services
 - Majordomo protocol

Security: there is none

- No security at all in ØMQ
 - assumes that you are on a private, internal network
- If you need to run "in the wild", use SSL:
 - bind ØMQ socket on localhost
 - run stunnel (with client cert verification)
- In progress: authentication layer
- dotCloud API is actually ZeroRPC, exposed through a HTTP/ZeroRPC gateway
- In progress: standardization of this gateway

Tracing (not published yet)

- Initial implementation during a hack day
 - bonus: displays live latency and request rates, using <http://projects.nuttnet.net/hummingbird/>
 - bonus: displays graphical call flow, using <http://raphaeljs.com/>
 - bonus: send exceptions to airbrake/sentry
- Refactoring in progress, to "untie" it from the dotCloud infrastructure and Open Source it

*How it works: all calls and responses are logged to a central place, along with a **trace_id** unique to each sequence of calls.*

Tracing: **trace_id**

- Each call has a **trace_id**
- The **trace_id** is propagated to subcalls
- The **trace_id** is bound to a local context (think thread local storage)
- When making a call:
 - if there is a local **trace_id**, use it
 - if there is none ("root call"), generate one (GUID)
- **trace_id** is passed in all calls and responses

Note: this is not (yet) in the github repository

Tracing: trace collection

- If a message (sent or received) has a **trace_id**, we send out the following things:
 - **trace_id**
 - call name (or, for return values, OK|ERR+exception)
 - current process name and hostname
 - timestamp

*Internal details: the collection is built on top of the standard **logging** module.*

Tracing: trace storage

- Traces are sent to a Redis key/value store
 - each **trace_id** is associated with a list of traces
 - we keep some per-service counters
 - Redis persistence is disabled
 - entries are given a TTL so they expire automatically
 - entries were initially JSON (for easy debugging)
 - ... then "compressed" with msgpack to save space
 - *approximately* 16 GB of traces per day

Internal details: the logging handler does not talk directly to Redis; it sends traces to a collector (which itself talks to Redis).

The problem with being synchronous

- Original implementation was synchronous
- Long-running calls blocked the server
- Workaround: multiple workers and a hub
- Wastes resources
- Does not work well for *very long* calls
 - Deployment and provisioning of new cluster nodes
 - Deployment and scaling of user apps

Note: this is not specific to ZeroRPC
(Preforking servers, threaded servers, WSGI...)

First shot at asynchronicity

- Send asynchronous events & setup callbacks
- "Please do **foo(42)** and send the result to this other place once you're done"
- We tried this. We failed.
 - distributed spaghetti code
 - trees falling in the forest with no one to hear them
- Might have worked better if we had...
 - better support in the library
 - better naming system
 - something to make sure that we don't lose calls (a kind of distributed FSM, maybe?)

Gevent to the rescue!

- Write synchronous code
(a.k.a. don't rewrite your services)
- Uses coroutines to achieve concurrency
- No fork, no threads, ~~no problems~~
- Monkey patch of the standard library
(to replace blocking calls with async versions)
- Achieve "unlimited" concurrency server-side

The version published on github uses gevent.

Show me the code!

<https://github.com/dotcloud/zerorpc-python>

\$ pip install git+git://github.com/dotcloud/zerorpc-python.git

Has: **zerorpc** module, **zerorpc-client** helper,
exception propagation, **gevent** integration

Doesn't have: tracing, naming, helpers for
PUB/SUB & PUSH/PULL, authentication

Questions?

Thanks!